

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
In [2]: titanic_train = pd.read_csv("train.csv", index_col = 0)
titanic_test = pd.read_csv("test.csv", index_col = 0)
titanic_train.head()
```

Out[2]:

	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
PassengerId											
1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

Analysis of the train dataset

```
In [3]: titanic_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 891 entries, 1 to 891
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    int64
1   Pclass      891 non-null    int64
2   Name        891 non-null    object
3   Sex         891 non-null    object
4   Age         714 non-null    float64
5   SibSp       891 non-null    int64
6   Parch       891 non-null    int64
7   Ticket      891 non-null    object
8   Fare        891 non-null    float64
9   Cabin       204 non-null    object
10  Embarked    889 non-null    object
dtypes: float64(2), int64(4), object(5)
memory usage: 83.5+ KB
```

```
In [4]: # Not all columns are needed for the analysis, drop the columns not needed.
column = ['Cabin', 'Ticket', 'Name']
titanic_train.drop(columns = column, inplace = True)
```

```
In [5]: titanic_train.isna().sum()
```

Out[5]:

Survived	0
Pclass	0
Sex	0
Age	177
SibSp	0
Parch	0
Fare	0
Embarked	2

dtype: int64

```
In [6]: #fill age column with median
titanic_train.Age.fillna(titanic_train.Age.median(), inplace = True)
```

```
In [7]: titanic_train.isna().sum()
```

Out[7]:

Survived	0
----------	---

```
Out[7]: Survived      0
Pclass      0
Sex          0
Age          0
SibSp        0
Parch        0
Fare         0
Embarked     2
dtype: int64
```

```
In [31]: titanic_train.dropna(inplace = True)
```

```
In [9]: titanic_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 889 entries, 1 to 891
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    889 non-null    int64
1   Pclass      889 non-null    int64
2   Sex         889 non-null    object
3   Age         889 non-null    float64
4   SibSp       889 non-null    int64
5   Parch       889 non-null    int64
6   Fare        889 non-null    float64
7   Embarked    889 non-null    object
dtypes: float64(2), int64(4), object(2)
memory usage: 62.5+ KB
```


```
In [10]: titanic_train.head()
```

```
Out[10]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
PassengerId								
1	0	3	male	22.0	1	0	7.2500	S
2	1	1	female	38.0	1	0	71.2833	C
3	1	3	female	26.0	0	0	7.9250	S
4	1	1	female	35.0	1	0	53.1000	S
5	0	3	male	35.0	0	0	8.0500	S

```
In [11]: #survivors based on sex
sns.countplot(data = titanic_train, x = 'Sex', hue = 'Survived')
```

```
Out[11]: <AxesSubplot:xlabel='Sex', ylabel='count'>
```



Convert the non numerical data to numerical data so as to predict the model

```
In [12]: #Split data into features and label (X and y)
X = titanic_train.drop('Survived', axis = 1)
y = titanic_train['Survived']
```

```
In [13]: # Transform the non numerical into numerical using sklearn
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer

cat_features = ['Sex', 'Embarked']
one_hot = OneHotEncoder()
transformer = ColumnTransformer([("one_hot", one_hot,
                                cat_features)],
                                remainder = "passthrough")

transformed_X = transformer.fit_transform(X)
transformed_X
```

```
Out[13]: array([[ 0.    ,  1.    ,  0.    , ...,  1.    ,  0.    ,  7.25   ],
 [ 1.    ,  0.    ,  1.    , ...,  1.    ,  0.    , 71.2833],
 [ 1.    ,  0.    ,  0.    , ...,  0.    ,  0.    ,  7.925   ],
 ...,
 [ 1.    ,  0.    ,  0.    , ...,  1.    ,  2.    , 23.45   ],
 [ 0.    ,  1.    ,  1.    , ...,  0.    ,  0.    , 30.     ],
 [ 0.    ,  1.    ,  0.    , ...,  0.    ,  0.    ,  7.75   ]])
```

```
In [14]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
np.random.seed(20)

#split data into train, test
X_train, X_test, y_train, y_test = train_test_split(transformed_X, y, test_size = 0.2)

clf = RandomForestClassifier(n_estimators = 100)
clf.fit(X_train, y_train)
```

```
Out[14]: RandomForestClassifier()
```

```
In [15]: clf.score(X_test, y_test)
```

```
Out[15]: 0.7752808988764045
```

```
In [16]: y_preds = clf.predict(X_test)
```

```
In [17]: pd.DataFrame({'Actual': y_test, 'Predicted': y_preds})
```

```
Out[17]:
```

	Actual	Predicted
PassengerId		
349	1	1
562	0	0
791	0	0
837	0	0
57	1	1
...
423	0	0
827	0	1
430	1	0
433	1	0
563	0	0


178 rows × 2 columns

```
In [18]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, precision_score, f1_score, r2_score
conf_mat = confusion_matrix(y_test, y_preds)

#Visualize confusion matrix using seaborn
def plot_mat(conf_mat):
```

```
fig, ax = plt.subplots(figsize = (3,3))
ax = sns.heatmap(conf_mat,
                  annot = True,
                  cbar = False)
plt.xlabel("Actual")
plt.ylabel("Predicted");

plot_mat(conf_mat)
```



```
In [19]: # evaluate the classifier
print('Classifier metrics on titanic data set')

print(f'Accuracy: {accuracy_score(y_test, y_preds)*100:.2f}%')
print(f'Precision : {precision_score(y_test, y_preds)}')
print(f'Recall: {recall_score(y_test, y_preds)}')
print(f'F1: {f1_score(y_test, y_preds)}')
```

```
Classifier metrics on titanic data set
Accuracy: 77.53%
Precision : 0.6891891891891891
Recall: 0.75
F1: 0.7183098591549296
```

The model prediction score is 77.53%

Using the model built, predict the test data

```
In [20]: titanic_test.head()
```

```
Out[20]:
```

	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
PassengerId										
892	3	Kelly, Mr. James	male	34.5	0	0	330911	7.8292	NaN	Q
893	3	Wilkes, Mrs. James (Ellen Needs)	female	47.0	1	0	363272	7.0000	NaN	S
894	2	Myles, Mr. Thomas Francis	male	62.0	0	0	240276	9.6875	NaN	Q
895	3	Wirz, Mr. Albert	male	27.0	0	0	315154	8.6625	NaN	S
896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	22.0	1	1	3101298	12.2875	NaN	S

```
In [21]: titanic_test.drop(columns = ['Ticket', 'Cabin', 'Name'], axis = 1, inplace = True)
titanic_test.Age.fillna(titanic_test.Age.median(), inplace = True)
titanic_test.fillna(0,inplace = True)
```

```
In [22]: titanic_test.head()
```

```
Out[22]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
PassengerId							
892	3	male	34.5	0	0	7.8292	Q
893	3	female	47.0	1	0	7.0000	S
894	2	male	62.0	0	0	9.6875	Q
895	3	male	27.0	0	0	8.6625	S
896	3	female	22.0	1	1	12.2875	S

```
In [23]: # convert non numerical to numerical
cate_features = ['Sex', 'Embarked']
one_hot = OneHotEncoder()
transformer = ColumnTransformer([('one_hot',
                                one_hot,
                                cate_features)],
                                remainder='passthrough')
transformed_data = transformer.fit_transform(titanic_test)
```

```
In [24]: predicted_test_data = (clf.predict(transformed_data))
```

```
In [25]: gender = pd.read_csv('gender_submission.csv')
gender.head()
```

```
Out[25]:
```

	PassengerId	Survived
0	892	0
1	893	1
2	894	0
3	895	0
4	896	1

```
In [26]: predicted_test_data
```

```
Out[26]: array([0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1,
1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1,
1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1,
1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1,
0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1,
0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0,
1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1,
0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1,
0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0,
1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0,
0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0,
1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1],
dtype=int64)
```

```
In [29]: #Load and save prediction
prediction = pd.DataFrame({'PassengerId': gender.PassengerId, 'Survived': predicted_test_data})
prediction.to_csv('submission.csv', index = False)
print('Your prediction was successfully saved')
```

Your prediction was successfully saved

```
In [30]: #check the saved prediction
test_pred = pd.read_csv('submission.csv', index_col = 0)
test_pred.head()
```

```
Out[30]:
```

	Survived
PassengerId	
892	0
893	0
894	0
895	1
896	1

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js